

# Text Representation Techniques

---

## Natural Language Processing

Daniel Ortiz Martínez, David Buchaca Prats

# Basic Definitions

- **Corpus:** a set of documents
- **Vocabulary:** a vocabulary is a set of words
- **Corpus vocabulary:** a set of words containing the atomic symbols used to represent the corpus
- **Token:** a substring from a document with *atomic* meaning (usually tokens refer to words)
- Vocabularies are usually extracted tokenizing a corpus

# Basic Text Representation

- In order to input text to a machine learning algorithm we need to convert the string representation to vectors
- The most basic way to encode text is a bag of words representation
  - A bag-of-words describes the occurrence of words within a text.
  - A bag of words representation involves:
    - A vocabulary of known words
    - A measure of the presence of known words (such as word counts)

# Basic Text Representation

- The vocabulary is usually stored as a dictionary (`Dict` or `OrderedDict`) that we will call `word_to_pos`
- Keys in `word_to_pos` are the words in the vocabulary
- Values in `word_to_pos` are the positions assigned to the words
- The bag of words feature vector  $\mathbf{x}$  for a document  $d$  is constructed using the counts of the words in  $d$
- Coordinate  $k$  in  $\mathbf{x}$  contains the number of times the  $k$ 'th word from `word_to_pos` appears in  $d$

# Basic Text Representation

```
word_to_pos = {'the':0, 'man':1, 'that':2, 'went':3, 'to':4, 'moon':5}
```

“The man that went to the moon”  $\rightarrow x = [1, 1, 1, 1, 1, 1]$

# Examples Bag of Words Representation

- Consider the corpus:

```
The cat sat on the mat
the cat and the dog sat on the mat
```

- Resulting in the following vocabulary:

```
word_to_pos = {'the':0, 'cat':1, 'sat':2, 'on':3, 'the':4, 'mat':5, 'and':6, 'dog':7}
```

- Bag of words representations:

“The cat that sat on the mat” [0, 1, 1, 1, 2, 1, 0, 0]

“the cat and the dog sat on the mat” [0, 1, 1, 1, 3, 1, 1, 1]

# Dictionaries for Bag of Words Representation: Standard Dict

```
normal_dict = {}  
normal_dict['1'] = "A"  
normal_dict['2'] = "B"  
normal_dict['3'] = "C"  
normal_dict['4'] = "D"  
normal_dict['5'] = "E"  
  
print("Printing normal dictionary:")  
for k,v, in normal_dict.items():  
    print("key : {0}, value: {1}".format(k,v))
```

```
Printing normal dictionary:  
key : 3, value: C  
key : 2, value: B  
key : 1, value: A  
key : 4, value: D  
key : 5, value: E
```

# Dictionaries for Bag of Words Representation: Ordered Dict

```
import collections

ordered_dict = collections.OrderedDict()
ordered_dict['1'] = "A"
ordered_dict['2'] = "B"
ordered_dict['3'] = "C"
ordered_dict['4'] = "D"
ordered_dict['5'] = "E"

print("Printing ordered dictionary:")
for k,v, in ordered_dict.items():
    print("key : {0}, value: {1}".format(k,v))
```

Printing normal dictionary:

```
key : 1, value: A
key : 2, value: B
key : 3, value: C
key : 4, value: D
key : 5, value: E
```



- How can we apply machine learning techniques when the input is a text description?
  - We need to transform strings to vectors
- Challenges when working with text:
  - The feature vector dimensionality can be huge
  - Words outside the vocabulary (such as misspelled words) might bring problems

# Creating Feature Vectors

- Given a corpus, how do we define a vocabulary?
  - We need to iterate over words, but raw data is not provided with words
- There are several decisions that impact vocabulary creation:
  - How do we generate tokens?
  - Do we need to clean tokens?
  - Do we create combinations of tokens?
  - Do we select combinations of tokens?

# Constructing a Vocabulary

- In order to build feature descriptors we need to create a `word_to_pos`
- `word_to_pos` depends on
  - How do we generate tokens from strings
  - How do we clean the tokens
- Many packages contain classes to create vectorizers with a `fit` method
- scikit-learn has the `CountVectorizer` class during fitting
  - Receives as input an iterable of strings
  - For each string the class finds the tokens (or words) in the string
  - Words are stored in `word_to_pos`

- `.fit(X)` learns a vocabulary from raw data `X`
- `.transform(x)` generates an array with the feature descriptor for `x`
- How should we store `.transform(x)`?
  - Numpy array
  - List
  - Pandas dataframe
  - **Scipy Compressed Sparse Row (CSR) matrix**

# High Dimensional Feature Vectors

- In the context of document descriptors feature vectors can contain millions of words
- Feature vectors are often sparse
- Storing such vectors using lists of Numpy arrays is very inefficient

# Compressed Sparse Row Matrix (CSR Matrix)

- Given an  $m \times n$  matrix, a CSR matrix is constructed from 3 arrays:
  - data**: contains the non-zero values of the matrix
  - ind\_col**: contains the column indices of the elements in the matrix
  - ind\_ptr**: contains  $m + 1$  pointers, the  $m$ 'th pointer stores the position of the first value in the **data** array that belongs to the  $m$ 'th row. The last element is equal to the length of the **data** array

```
X = np.array([[0, 5, 7, 6, 0, 0, 0, 0, 0, 0],
              [0, 0, 0, 0, 0, 0, 0, 0, 0, 1],
              [7, 0, 4, 9, 0, 0, 0, 0, 0, 0]])

data      = [5, 7, 6, 1, 7, 4, 9]
ind_col   = [1, 2, 3, 9, 0, 2, 3]
ind_ptr   = [0, 3, 4, 7]

X_csr = sp.csr_matrix((data, ind_col, ind_ptr))
X_csr.toarray()

array([[0, 5, 7, 6, 0, 0, 0, 0, 0, 0],
       [0, 0, 0, 0, 0, 0, 0, 0, 0, 1],
       [7, 0, 4, 9, 0, 0, 0, 0, 0, 0]])
```

# Transforming an Iterable of Strings to a Sparse Matrix

- How can we transform a sequence of strings to a sparse matrix?
- One approach would be:
  1. Iterate over the data to create a vocabulary
  2. Iterate over the data to generate the feature vectors for the elements in the vocabulary
- Can we improve this?

# Generating a Feature Descriptor for a List of Documents

```
docs = [['hello', 'world', 'hello'], ['goodbye', 'cruel', 'teacher' ]]  
  
def prepare_word_counts_with_dict(docs, verbose=False):  
    ind_ptr = [0]  
    ind_col = []  
    data = []  
    vocabulary = {}  
  
    for m, doc in enumerate(docs):  
        # TO-BE-DONE: Use an auxiliary dictionary to keep track of counts  
        #  
        #  
        #  
        #  
        #  
        #  
        #  
        #  
        #  
        return (data, ind_col, ind_ptr)  
  
sp.csr_matrix(prepare_word_counts_with_dict(docs)).toarray()  
  
array([[2, 1, 0, 0, 0],  
       [0, 0, 1, 1, 1]])
```



# Generating a Feature Descriptor for a List of Documents

```
docs = [['hello', 'world', 'hello'], ['goodbye', 'cruel', 'teacher']]

def prepare_word_counts_with_dict(docs, verbose=False):
    ind_ptr = [0]
    ind_col = []
    data = []
    vocabulary = {}

    for m, doc in enumerate(docs):
        word_ind_counter = collections.defaultdict(int)
        for word in doc:
            vocabulary.setdefault(word, len(vocabulary))
            word_ind_counter[word] += 1

        data.extend(word_ind_counter.values())
        ind_ptr.append(ind_ptr[-1] + len(word_ind_counter))
        ind_col.extend([vocabulary[w] for w in word_ind_counter.keys()])

    return (data, ind_col, ind_ptr)

sp.csr_matrix(prepare_word_counts_with_dict(docs)).toarray()

array([[2, 1, 0, 0, 0],
       [0, 0, 1, 1, 1]])
```

# Generating a Feature Descriptor for a List of Documents


```
docs = [['hello', 'world', 'hello'], ['goodbye', 'cruel', 'teacher' ]]  
  
ind_ptr = [0]  
ind_col = []  
data = []  
vocabulary = {}  
  
for d in docs:  
    # TO-BE-DONE: Implementation WITHOUT auxiliary dictionary  
    #  
    #  
    #  
    #  
  
sp.csr_matrix((data, ind_col, ind_ptr)).toarray()  
  
array([[2, 1, 0, 0, 0],  
       [0, 0, 1, 1, 1]])
```

# Generating a Feature Descriptor for a List of Documents


```
docs = [['hello', 'world', 'hello'], ['goodbye', 'cruel', 'teacher' ]]  
  
ind_ptr = [0]  
ind_col = []  
data = []  
vocabulary = {}  
  
for d in docs:  
    for term in d:  
        index = vocabulary.setdefault(term, len(vocabulary))  
        ind_col.append(index)  
        data.append(1)  
    ind_ptr.append(len(ind_col))  
  
sp.csr_matrix((data, ind_col, ind_ptr)).toarray()  
  
array([[2, 1, 0, 0, 0],  
       [0, 0, 1, 1, 1]])
```

# Different Ways to Cast a Learning Task


- Typical ML pipeline




- Composable/composite models



# Examples



```
count_vectorizer = sklearn.feature_extraction.text.CountVectorizer()  
logistic = sklearn.linear_model.LogisticRegression(C=0.1)  
  
X_train_feature_vec = count_vectorizer.fit_transform(X_train)  
logistic.fit(X_train_feature_vec, y_train)
```

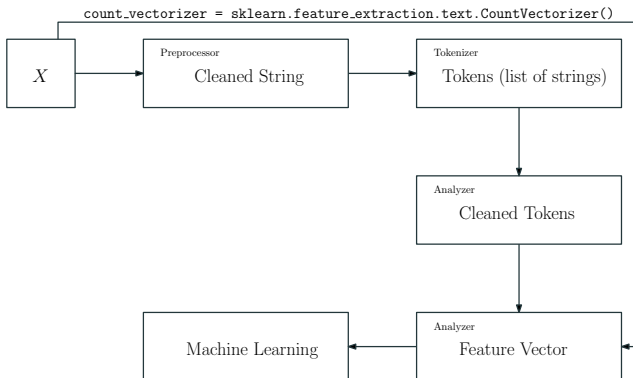


```
count_vectorizer = sklearn.feature_extraction.text.CountVectorizer()  
logistic = sklearn.linear_model.LogisticRegression(C=0.1)  
  
model_pipe = sklearn.pipeline.Pipeline([("countvectorizer", count_vectorizer),  
                                         ("logisticregression", logistic)])  
  
model_pipe.fit(X_train, y_train)
```

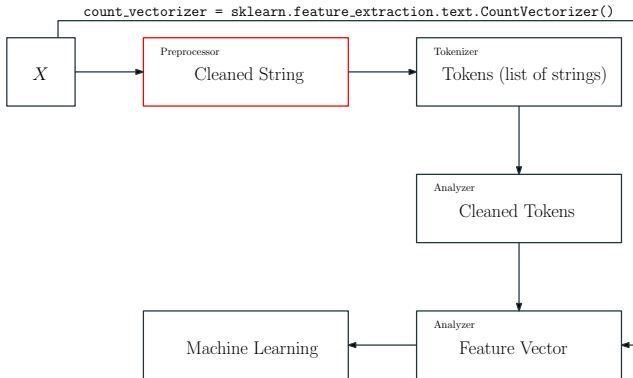
# Pipelines or Composite Models

- The purpose of the pipeline is to assemble a composite model which consists on several steps that can be cross-validated together
- Pipelines allow practitioners to easily compose and validate decisions made during training, instead of relying on pre-processing steps with *hand-crafted* decisions
- Some of this decisions might include:
  - Type of tokenizer
  - Removing stop words
  - Using only words or bigram, trigrams etc..
  - Regularization parameters

- A `CountVectorizer` is one of the most straight forward methods to generate descriptors for documents



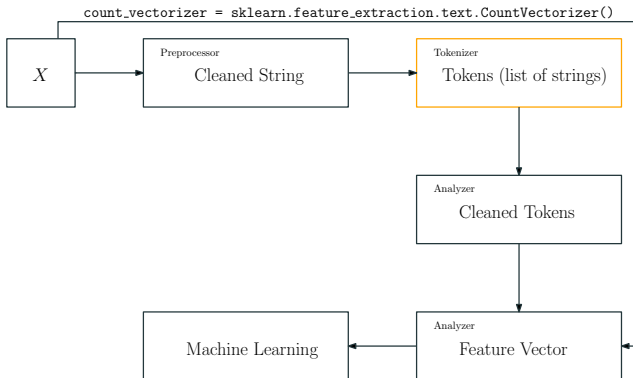
# CountVectorizer: Preprocessor



```
count_vectorizer = sklearn.feature_extraction.text.CountVectorizer()  
count_vectorizer.fit(X_train)  
CountVectorizer(analyzer='word', binary=False, decode_error='strict',  
                dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',  
                lowercase=True, max_df=1.0, max_features=None, min_df=1,  
                ngram_range=(1, 1), preprocessor=None, stop_words=None,  
                strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',  
                tokenizer=None, vocabulary=None)
```



# CountVectorizer: Tokenizer



```
count_vectorizer = sklearn.feature_extraction.text.CountVectorizer()  
count_vectorizer.fit(X_train)  
CountVectorizer(analyzer='word', binary=False, decode_error='strict',  
                dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',  
                lowercase=True, max_df=1.0, max_features=None, min_df=1,  
                ngram_range=(1, 1), preprocessor=None, stop_words=None,  
                strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',  
                tokenizer=None, vocabulary=None)
```

# Tokenizer limitations

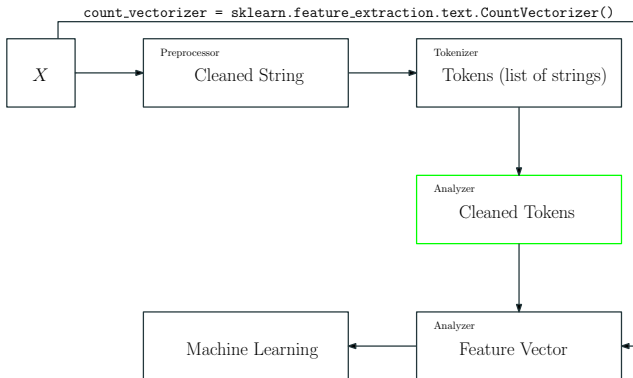
- Single letter words such as “I” are ignored

```
re.findall(r'(?u)\b\w\w+\b', "I can't wait to go there!")  
['can', 'wait', 'to', 'go', 'there']
```

- Expressions such as “can't” are modified and might even change meaning

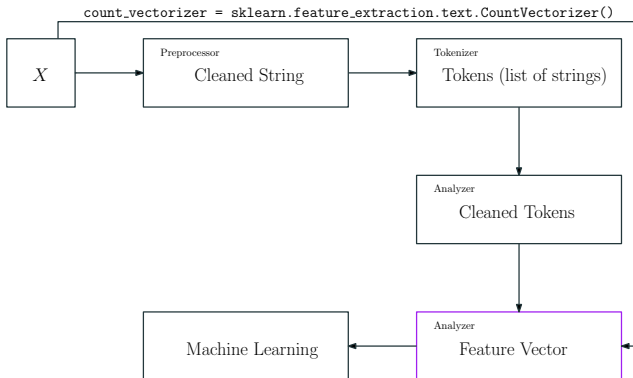
```
re.findall(r"\w+\'\w+", "I can't wait but I won't go there")  
["can't", "won't"]
```

# CountVectorizer: Analyzer



```
count_vectorizer = sklearn.feature_extraction.text.CountVectorizer()  
count_vectorizer.fit(X_train)  
CountVectorizer(analyzer='word', binary=False, decode_error='strict',  
                dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',  
                lowercase=True, max_df=1.0, max_features=None, min_df=1,  
                ngram_range=(1, 1), preprocessor=None, stop_words=None,  
                strip_accents=None, token_pattern='(?u)\\b\\w+\\b',  
                tokenizer=None, vocabulary=None)
```

# CountVectorizer: Analyzer



```
count_vectorizer = sklearn.feature_extraction.text.CountVectorizer()  
count_vectorizer.fit(X_train)  
CountVectorizer(analyzer='word', binary=False, decode_error='strict',  
                dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',  
                lowercase=True, max_df=1.0, max_features=None, min_df=1,  
                ngram_range=(1, 1), preprocessor=None, stop_words=None,  
                strip_accents=None, token_pattern='(?u)\\b\\w+\\b',  
                tokenizer=None, vocabulary=None)
```

## Summary: Customizing Vectorizer Classes

- **Preprocessor:** a callable that takes an entire document as input (as a single string), and returns a possibly transformed version of the document, still as an entire string. This can be used to remove HTML tags, lowercase the entire document, etc.
- **Tokenizer:** a callable that takes the output from the preprocessor and splits it into tokens, then returns a list of these
- **Analyzer:** a callable that replaces the preprocessor and tokenizer. The default analyzers all call the preprocessor and tokenizer, but custom analyzers will skip this. N-gram extraction and stop word filtering take place at the analyzer level, so a custom analyzer may have to reproduce these steps

# Why Is Important to Tune Vectorizers

- Many times vocabulary can be so rare that is not worth storing it
- One good example would be numbers, alphanumeric codes, etc.
- These words can be categorized, using special symbols to represent them

# Word Transformations: Stemming

- Stemming consist on removing the suffixes or prefixes used in word
- The returned string from a stemmer might not be a valid word from the language
- Example:

```
Stem(saw) = saw  
Stem(destabilize) = destabil
```

# Word Transformations: Lemmatization

- Lemmatization consist on properly use of a vocabulary and morphological analysis of words, aiming to remove inflectional endings only with the goal of returning any word to a set of base (or dictionary form) words
- The returned string from a lemmatizer should be a valid word from the language
- Example:

```
Lemmatize(saw) = see
```

```
Lemmatize(destabilize) = destabilize
```