

# BK-Trees: Efficient Retrieval of Similar Strings

---

**Natural Language Processing**

Daniel Ortiz Martínez, David Buchaca Prats

# Building a Basic Spellchecker

- Based on a notion that we have already seen, edit distance, we can build a spellchecker
- Let us consider that we want to correct misspelled words (i.e. words that are not in the vocabulary)
- Base algorithm: Let  $x$  be a sentence and  $V$  the vocabulary

```
for w in x
  if w not in V
    Find the closest words to w # (candidate search)
    Evaluate each of the candidate words # (candidate evaluation)
    Return the most probable candidate
```

# Finding the Closest Items to a Query

- This is an extremely relevant problem for many data science and machine learning problems
- Scikit-learn implements the kdtree for dense vectors
- What if we have strings?

# Finding the Closest Items to a Query

- Efficient search of similar values in a dataset is a very challenging problem. In particular, computing distances between a word and a huge vocabulary can be computationally expensive
- Let  $w$  be a string that is out of the vocabulary
- Let us consider  $W_k(w; X) = \{w | w \in V, d(w, w_j) < k\}$
- Finding  $W_k(w; X)$  can be done in two different ways:
  1. compute  $d(w, w_j)$  for all  $w_j$  keeping elements at distance at most  $k$
  2. Use a data structure to avoid computing  $d(w, w_j)$  for all  $w_j$  in  $X$

# Tree Intuition

- We can build a tree to do efficient search of similar words. This will allow us to prune a lot of the search space, with the objective of avoiding many distance computations on a big part of the vocabulary
- Example: consider  $w = \text{pleistation}$   
ana  $\rightarrow d(\text{pleistation}, \text{ana})$   
playstation  $\rightarrow d(\text{pleistation}, \text{playstation})$   
house  $\rightarrow d(\text{pleistation}, \text{house})$
- If “pleistation” has 11 characters and we want candidates at most at distance  $k = 3$ , is there any need to compute  $d(\text{pleistation}, \text{ana})$  ?  
Ana has 3 characters!

# BK-Tree: Construction Algorithm


1. Select any word from the vocabulary and use it as the root node
2. Keep adding words until all vocabulary is in the tree
  - 2.1 Each time we add a word the distance between the word and the root node is computed, let us assume this distance is  $d$
  - 2.2 If no node from the root node is at distance  $d$  we add a new leaf as a descendant of the root node with edge value equal to  $d$
  - 2.3 If there exist another node at distance  $d$  then we repeat this process redefining the root node as the node that produced the collision

# BK-Tree: Construction Example


- Let us consider the data [book, books, cake, boo, cape, cart, boon, cook]
- Insert **book** (which becomes root node)



- Insert **books**: compute  $d(\text{book}, \text{books})=1$




- Insert **cake**: compute  $d(\text{book}, \text{cake})=4$



# BK-Tree: Construction Example


- Insert **boo**: compute  $d(\text{book}, \text{boo})=1$ 
  - The BK-tree has to respect that every node have all children with different distances, since there is already a word at the same edit distance 1 we go to the branch of words at distance 1
  - If there is a collision (like we have now) the new word must become a children of the collided word. In this case, a children of “book”
  - The new distance from “books” to “boo” is 2






# BK-Tree: Construction Example

- Insert **cape**: compute  $d(\text{book}, \text{cape})=4$ 
  - Collision! There is already cake at distance 4 from “book”
  - Root node is now “cake”
  - Root=cake: compute  $d(\text{cake}, \text{cape})=1$
  - There is no descendant from “cake” at distance 1  $\rightarrow$  we can add it




# BK-Tree: Construction Example

- Insert **cart**: compute  $d(\text{book}, \text{cart})=4$ 
  - Collision! There is already cake at distance 4 from “book”
  - Root node is now “cake”
  - Root=cake: compute  $d(\text{cake}, \text{cart})=2$
  - There is no descendant from “cake” at distance 2  $\rightarrow$  we can add it




# BK-Tree: Construction Example

- Insert **boon**: compute  $d(\text{book}, \text{cart})=4$ 
  - Collision! There is already cake at distance 4 from “book”
  - Root node is now “books”
  - Root=books: compute  $d(\text{books}, \text{boon})=2$
  - Collision! There is already “boo” at distance 2 from “boon”
  - Root=books: compute  $d(\text{boo}, \text{boon})=1$
  - There is no descendant from “boo” at distance 1  $\rightarrow$  we can add it




# BK-Tree: Construction Example

- Insert **cook**: compute  $d(\text{book}, \text{cook})=1$ 
  - Collision! There is already cake at distance 1 from “book”
  - Root node is now “books”
  - Root=books: compute  $d(\text{books}, \text{cook})=2$
  - Collision! There is already “boo” at distance 2 from “cook”
  - Root=books: compute  $d(\text{boo}, \text{cook})=2$
  - There is no descendant from “boo” at distance 2  $\rightarrow$  we can add it



# BK-Tree: Storage in Memory


- Let us consider the following tree:



- We can use tuples to represent the tree in memory:
  - The first element is the word assigned to the node
  - The second element is the subtree that spawns from that node
  - A subtree can be represented as a Dict[Int, Tuple]
  - Keys are the distances to the root node
  - Values are tuples which represent subtrees
- For the previous example:


```
('book', {1: ('books', {}), 4: ('cake', {})})
```

# BK-Tree: Storage in Memory



```
('book',  
{1: ('books', {2: ('boo', {})}),  
4: ('cake', {}))}
```


# BK-Tree: Storage in Memory



```
('book',  
{1: ('books', {2: ('boo', {1: ('boon', {}), 2: ('cook', {}))})}),  
4: ('cake', {1: ('cape', {}), 2: ('cart', {}))})})
```

# Searching in a BK-Tree


- Problem: search all words that appear at distance less or equal than a tolerance  $T$  from a query word  $q$
- Bad solution: compute all edit distances between  $q$  and  $w$  for  $w$  in the vocabulary
- Key idea: visit all words  $w$  that are at distance  $[d(w, q) - T, d(w, q) + T]$






# BK-Tree: Searching Example

- Let us consider  $q=\text{cage}$ ,  $T=1$ ,  $\text{candidates}=[]$ ,  $\text{search}=[\text{book}]$
- Select candidate “book” from  $\text{search}=[\text{book}]$ 
  - $d(\text{book}, \text{cage}) = 4 \rightarrow \text{candidates}$  is not updated
  - Crawl all children of “book” at distance  $I=[4-1, 4+1]=[3,5]$
  - Only node cake is connected to book and with distance in  $I=[3,5]$
  - $\text{search} = [\text{book}, \text{cake}] \setminus \text{book} = [\text{cake}]$




# BK-Tree: Searching Example

- Let us consider  $q=\text{cake}$ ,  $T=1$ ,  $\text{candidates}=[]$
- Select candidate “cake” from  $\text{search}=[\text{cake}]$ 
  - $d(\text{cake}, \text{cake}) = 1 \rightarrow \text{candidates} += [\text{cake}]$
  - Crawl all children of “cake” at distance  $I=[1-1, 1+1]=[0, 2]$
  - There are only 2 possible nodes,  $\text{search} = [\text{cape}, \text{cart}]$




# BK-Tree: Searching Example

- Let us consider  $q=\text{cape}$ ,  $T=1$ , candidates= $[\text{cake}]$
- Select candidate “cape” from search= $[\text{cape}, \text{cart}]$ 
  - $d(\text{cape}, \text{cape}) = 1 \rightarrow \text{candidates} += [\text{cape}]$
  - Crawl all children of “cape” at distance  $I=[1-1, 1+1]=[0, 2]$
  - “cape” has no children
  - $\text{search} = [\text{cape}, \text{cart}] \setminus \text{cape} = [\text{cart}]$



# BK-Tree: Searching Example


- Let us consider  $q=\text{cage}$ ,  $T=1$ ,  $\text{candidates}=[\text{cake}, \text{cape}]$
- Select candidate “cape” from  $\text{search}=[\text{cart}]$ 
  - $d(\text{cart}, \text{cage}) = 2 \rightarrow \text{candidates}$  is not updated
  - Crawl all children of “cart” at distance  $I=[2-1, 2+1]=[1, 3]$
  - “cart” has no children
  - $\text{search} = [\text{cart}] \setminus \text{cart} = [] \rightarrow$  Search space is empty, stop search



- The resulting set of possible candidates at distance 1 are:  $[\text{cake}, \text{cape}]$

# BK-Tree: Searching Example

- To sum up:
  - Start conditions:  $q=\text{cage}$ ,  $T=1$ ,  $\text{candidates}=[]$ ,  $\text{search}=[\text{book}]$
  - Result: the set of possible candidates at distance 1 are  $[\text{cake}, \text{cape}]$
- Observation: we ended up computing 4 edit distances yet we have 8 nodes



# BK-Tree Speedup

- In the case that the search space is drastically pruned, the speedup can be massive:

```
word = "anthropomorphologically"  
max_dist = 2  
sort_candidates=False  
%timeit candidates_ext = get_candidates_exhaustive(word,max_dist,words)  
245 ms  $\pm$  8.68 ms per loop (mean  $\pm$  std. dev. of 7 runs, 1 loop each)
```

```
candidates_ext = get_candidates_exhaustive(word,max_dist,words)  
candidates_ext  
[(1, 'anthropomorphological'), (1, 'anthropomorphologically')]
```

```
word = "anthropomorphologically"  
%timeit candidates_ext = t.query(word, 2)  
123  $\mu$ s  $\pm$  805 ns per loop (mean  $\pm$  std. dev. of 7 runs, 10,000 loops each)
```

```
candidates_ext = t.query(word, 2)  
candidates_ext  
[(1, 'anthropomorphological'), (1, 'anthropomorphologically')]
```

# BK-Tree Speedup

- If the pruned search space still contains a huge amount of words the speedup might not be that huge:

```
word = "astrologi"  
max_dist = 2  
sort_candidates=False  
%timeit candidates_ext = get_candidates_exhaustive(word, max_dist, words)
```

221 ms  $\pm$  14.7 ms per loop (mean  $\pm$  std. dev. of 7 runs, 1 loop each)

```
word = "astrologi"  
%timeit t.query(word, 2)
```

99.6 ms  $\pm$  724  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 10 loops each)