

Practice 14 R: GARCH models

1. White but not IID noise residuals.

Let us analyze the daily returns of the S&P500:

- `library(MASS)`
- `data(SP500)`
- `plot(SP500, type="l")`
- `acf(SP500)`
- `Box.test(SP500, type="Ljung-Box")`

Clearly we have a white noise.

We can compute the mean:

- `m<-mean(SP500)`
- `x<-SP500-m`
- `plot(x, type="l")`
- `acf(x)`

Of course, we are under the same situation. Note that the series x is not stationary in variance.

But do the following:

- `y<-x^2`
- `plot(y, type="l")`
- `acf(y)`
- `Box.test(y, type="Ljung-Box")`

What is happening? Note that x is a white noise but y not.

We need a new theory to give structure to this type of processes.

2. Generating ARCH and GARCH models.

- `a0<-0.25`
- `a1<-0.5`
- `x<-numeric(9999)`
- `x0<-rnorm(1)`
- `for (i in 1:9999){`
- `x0<-rnorm(1,sd=sqrt(a0+a1*x0^2))`
- `x[i]<-x0`
- `}`
- `plot(x, type="l")`

An alternative system for a GARCH(1,1) is

- `n<-10000`
- `a0<-0.1`
- `a1<-0.4`
- `b1<-0.2`
- `w<-rnorm(n)`
- `y<-rep(0,n)`
- `h<-rep(0,n)`
- `for(i in 2:n){`
- `h[i]<-a0+a1*(y[i-1]^2)+b1*h[i-1]`
- `y[i]<-w[i]*sqrt(h[i])`
- `}`
- `plot(y, type="l")`
- `acf(y)`
- `acf(y^2)`

3. Do the same with an ARCH(2).

4. Fitting GARCH models

Now we adjust a GARCH model to SP500 data.

```
library(MASS)
library(tseries)
data(SP500)
plot(SP500, type="l")
x<-SP500-mean(SP500)
plot(x,type="l")
sp<-garch(x, trace=F)
res<-sp$res[-1]
plot(res,type="l")
acf(res)
acf(res^2)
confint(sp)
garch(x,order=c(1,1), trace=F)
```

The instruction “confint” gives the confidence intervals of the parameters of the estimated GARCH model. The last instruction gives the estimated parameters.